

CIFAR-10 Image Classification

Objective

Build and train deep learning models using TensorFlow or PyTorch. Apply best practices in data handling, model architecture, training loops, and evaluation.

Part 1 — Core Coding Task

Build and train a neural network to classify images from the CIFAR-10 dataset.

Requirements

1. **Data Pipeline**
 - Load CIFAR-10 dataset
 - Apply standard preprocessing (normalization, augmentations, batching)
 - Demonstrate efficient data loading (tf.data or torch.utils.data.DataLoader)
2. **Model Architecture**
 - Implement a small CNN **from scratch** (not a pre-trained model), including:
 - Convolutional layers
 - BatchNorm
 - Dropout
 - A final classification head
3. **Training**
 - Implement a clear training loop:
 - Loss computation
 - Accuracy metrics
 - Learning rate scheduling
 - Save checkpoints
4. **Evaluation**
 - Plot loss & accuracy curves
 - Compute test accuracy
 - Show sample misclassified images
5. **Reproducibility**
 - Set random seeds
 - Provide environment/requirements file

Deliverables

- Source code (Python scripts or notebook)

- README explaining:
 - Project structure
 - Training setup
 - Model design decisions
 - How to run the code

Part 2 — Model Improvement Challenge

After the baseline model is working, implement **one** of the following improvements:

Option A — Add a Residual Block

Create a simplified version of a ResNet block and integrate it into your CNN.

Option B — Add a Custom Callback / Training Utility

Example:

- TensorBoard (TF)
- Custom PyTorch callback or logger
- Early stopping
- Gradient clipping

Option C — Hyperparameter Tuning

Perform tuning on:

- Learning rate
- Batch size
- Optimizer
- Dropout rate

Provide results comparing baseline vs improved model.

Optional Extensions

- Implement mixed-precision training
- Use distributed training (Multi-GPU / TPU / DDP)
- Deploy the trained model with FastAPI, Flask, or TensorFlow Serving
- Convert model to ONNX and perform inference
- Add unit tests for model components